

Cloud Computing Lecture Notes

Distributed Computing/Systems

Definition:

Distributed computing refers to a system where computing resources are distributed across multiple locations rather than being centralized in a single system. This enables task distribution and efficient resource utilization.

Why Use Distributed Systems?

- **Scalability Issues:** Traditional computing faces bottlenecks due to hardware limitations, whereas distributed systems allow for hardware scaling.
- **Connected Devices:** In a networked system, connected devices communicate, but this does not necessarily make them distributed.
- **IoT (Internet of Things):** IoT is one of the largest examples of distributed computing.
- **Multi-layered System Design:** Distributed computing enables systems to function in multiple layers, with each layer acting as a distributed entity.
- **User Perspective:** Although the system consists of multiple machines, distributed computing presents a unified system to users.

Parallel Computing

Parallel computing involves executing multiple processes simultaneously to enhance speed and efficiency.

Key Aspects of Parallel Computing:

- **Distribution & Speed:** Distributing tasks does not always guarantee faster execution; efficiency depends on the nature of the task and system design.
- **Core Objective:** The primary goal is to **speed up** processing by leveraging parallel programming and hardware optimizations.
- **Use Cases:** Parallel computing is widely used in:

- Vector processing
 - Image processing
 - Matrix multiplication
- **Limitations:**
 - Not all applications can be parallelized.
 - Some components of code can be executed in parallel, while others may not.
 - Specific programming languages are required for parallel computing.
- **Computing Infrastructure:**
 - Dedicated machines are needed for parallel computing.
 - **Clusters:** A collection of similar types of systems working together.

Building Distributed Systems

Key Components:

1. **Shared Folders:** Simply sharing files does not make a system truly distributed. Proper access control and system design are needed.
2. **Middleware:** The software layer that facilitates communication between different system components.

Diagram in Notes:

A simple representation of middleware connecting multiple system components. The figure depicts three components (labeled 1, 2, 3) connected via middleware, highlighting its role in integrating distributed elements.

Role of Middleware:

- A) Acts as a bridge between applications and networked devices.
- B) Provides a unified interface for interacting with connected components.

Clusters vs. Grids

Clusters:

Clusters consist of multiple machines with similar hardware and operating systems, working together in a **symmetric** manner.

- **Key Features:**
 - Middleware abstracts hardware, OS, and software details from users.
 - All machines have the **same type of hardware and OS**.
 - Parallel computing is implemented efficiently in clusters.
 - Middleware in clusters has fewer complexities.
 - Clusters typically exist in **LAN (Local Area Networks)**, meaning all machines are geographically close and in the same environment.

Grids:

Grids, unlike clusters, consist of **heterogeneous** systems that may have different hardware, OS, and configurations. They are **geographically distributed** and operate in an **asymmetric** manner.

- **Key Features:**
 - Middleware in grids has to handle diverse configurations, leading to potential **performance degradation**.
 - Used for solving large-scale computational problems.
 - Systems join grids for incentives, such as **resource sharing** or **financial compensation**.
 - Users can sell extra computational resources within a grid.
 - Example: **Blockchain technology** operates in a grid-like manner.
 - **Drawbacks:** Security concerns due to distributed nature.

Comparison of Clusters & Grids:

- Organizations **do not** allow their internal systems to be part of a grid due to security concerns.
- Individual users, however, can participate in grid computing.
- **Example:** Torrents are a similar concept but not an exact representation of grids.
- **Clusters & Grids together form the foundation of cloud computing.**

- They also contribute to **High-Performance Computing (HPC)**.

Cloud Computing

Definition:

Cloud computing is a form of **distributed computing** that provides on-demand computing resources over a network.

Characteristics of Cloud Computing:

- **No Operator Wait Time:** Users don't need to wait for an operator to grant access.
- **Ubiquitous Access:** Cloud services can be accessed from anywhere with an internet connection.
- **Utility Computing Model:** Users pay only for what they consume.
- **Existing Infrastructure:** Many foundational blocks of cloud computing are already present in **Data Centers (DCs)**.

Cloud Computing & HPC (High-Performance Computing):

- **HPC as a Cloud Service:** Cloud providers like **AWS** offer **HPC as a service**.
- **Cluster on Demand:** Users can request cluster-based resources directly from cloud providers.
- **Kubernetes Integration:** Cloud computing allows users to build their own **Kubernetes clusters** or use Kubernetes as a managed service.
- **Autoscaling:** Cloud computing dynamically adjusts resources based on demand.
- **Complex Workloads:** Supports operations like **Machine Learning (ML)** and **big data analytics**.

Supercomputing in the Cloud:

- **Top500.org:** Lists the world's top supercomputers.
- **Supercomputers:** Consist of **millions of cores**, forming clusters of multiple machines.
- **Difference from Cloud Computing:** Unlike the cloud, supercomputers grant access only to **specific users** requiring HPC.

Middleware in Cloud Computing:

- **Resource Management:** Middleware is responsible for managing resources in the cloud.
- **User Abstraction:** End users do not need to know where the cloud system's infrastructure is physically located.

Key Desired Attributes of Cloud Computing:

- **Backend Abstraction:** Cloud systems hide backend configurations and resource management from users.
- **Privacy & Compliance:** Some details, such as the **physical location of machines**, must be disclosed due to legal and compliance requirements.

Cloud Computing Overview

- If certain **essential characteristics** are present in distributed systems, they qualify as **cloud computing**.
- The **NIST (National Institute of Standards and Technology)** defines **cloud computing** as an extension of **distributed systems**.

Essential Characteristics of Cloud Computing

These characteristics distinguish **cloud computing** from traditional distributed systems.

1. On-Demand Self-Service

- Users can **configure and deploy services** independently, without requiring human intervention.
- No need to wait for an operator; services are **already deployed** and **provided via a web interface**.

2. Broad Network Access

- Cloud services establish a **fast** connection between servers and users.

- Since computing is performed **on the operator's side**, the results are delivered quickly to users.
- Users interact with the system via **APIs (Application Programming Interfaces)** or **web interfaces** without needing direct access to computing resources.

3. Resource Pooling

- The **resources of a cloud infrastructure** should be available to **all users**.
- Cloud providers combine **resources** into a shared pool.
- Users can utilize any resource as long as it is available, and once released, the resource **returns to the pool** for reallocation.

4. Rapid Elasticity

- Cloud computing **scales up or down** dynamically.
- Workloads can **expand or contract** based on demand.
- Scaling is **fast and automatic**, ensuring efficient resource utilization.
- Example: **AWS EC2 (Elastic Compute Cloud)** provides rapid elasticity.

5. Measured Service

- Cloud providers implement **pay-as-you-go** billing models.
- **Usage statistics** help with **future planning, security (e.g., anomaly detection), and billing transparency**.

Common Characteristics of Distributed & Cloud Systems

Cloud computing shares several **common characteristics** with general **distributed systems**.

1. Resource Pooling

- Present in **all distributed systems**, including **grid computing and clusters**.

2. Measured Services

- Found in **most distributed systems**.

- **Middleware** is responsible for **resource pooling**.
- Examples: **Hadoop clusters, Apache Spark**.

3. Essential vs. Common Characteristics

- The presence of common characteristics **alone does not** guarantee a system qualifies as **cloud computing**.
- Without these **common characteristics**, achieving the **essential characteristics** becomes extremely difficult.
- **Cloud construction without common characteristics is theoretically possible but practically impossible**.

4. Massive Scale

- A vast amount of **computing resources** is combined to form **large-scale cloud systems**.
- Some clouds operate on a **massive scale**, while others do not.
- **Private clouds for organizations are not massive**.
- **Public cloud market dominance: 95-96% of cloud infrastructure** belongs to a few major companies like:
 - **AWS (Amazon Web Services)**
 - **GCP (Google Cloud Platform)**
 - **Azure (Microsoft)**
 - **IBM Cloud**

5. Homogeneity

- Cloud systems prefer **symmetric resources** (same type of OS, servers, configurations, etc.).
- Although different resources can be used, **homogeneity** simplifies **management and performance**.
- **Asymmetric resources** lead to **performance degradation** due to compatibility challenges.

6. Virtualization

- Virtualization allows **creation of virtual objects** from physical resources, such as:

- **Virtual Memory**
- **Virtual Storage**
- **Virtual Operating Systems**
- **Virtualization alone is not cloud computing**, but it is a key enabler.
- **Benefits of Virtualization:**
 - Supports **resource pooling**.
 - Provides **elasticity**.
 - **Eases workload migration** (moving virtual machines is easier than moving physical machines).
- **All cloud services are virtualized.**

7. Resilient Computing

- Cloud resilience ensures system availability despite **failures, redundancies, and faults**.
- Includes **backup, fault tolerance, early failure detection, and mitigation strategies**.
- **Resilience is relative:**
 - **Massive-scale clouds** offer **higher resilience**.
 - **Smaller/private clouds** may have **lower resilience**.

8. Service Orientation

- Cloud resources are provided as **web services**.
- Services include:
 - **Computing**
 - **Storage**
 - **Networking**
- **SOA (Service-Oriented Architecture)**: A design approach used for cloud services.

9. Advanced Security

- **Security breaches impact a large number of users** in cloud environments.
- Requires **significant investment** in **cloud security**.
- Security is **relatively easier** to maintain at a **large scale** than at a **small scale**.

10. Low-Cost Software

- Large-scale cloud usage **reduces software costs** due to **economies of scale**.
- **High production & demand** allow software costs to decrease **relatively**.

11. Geographical Distribution

- Cloud infrastructure is **distributed across multiple locations** to:
 - Improve **fault tolerance**.
 - Enhance **resilience, efficiency, and speed**.
 - Ensure **scalability and privacy compliance**.
- **CDN (Content Delivery Network)**: A **distributed network of cloud instances** that **store and distribute content** to users efficiently.
- **Some clouds are not geographically distributed** but offer local compliance options.

Cloud Service Model

Cloud service models define:

1. **How users access services**
2. **How operators manage resources**
3. **Responsibility division between users and cloud service providers (CSPs)**

Responsibility Model

- **User vs. Cloud Service Provider (CSP)**
 - Defines **who manages what** in cloud-based systems.

(Diagram Present in Notes)

- A diagram represents a **responsibility model** where different layers of the cloud stack (Infrastructure, Platform, and Software) are managed either by the **user** or the **CSP**.

1. Infrastructure as a Service (IaaS)

- Example: **AWS EC2**
- Users manage:
 - **Applications**
 - **Operating Systems**
 - **Virtual Machines (VMs)**
 - **Servers**
 - **Networking**
- The **cloud provider** manages the underlying **hardware**.

(Diagram Present in Notes)

- A diagram illustrating **IaaS** shows how users handle **VMs, servers, and networking**, while the cloud provider manages the **infrastructure**.

2. Platform as a Service (PaaS)

- Users interact with **applications** but do not manage the **underlying platform**.
- Example: **Kubernetes**
- The **operator** manages the **platform**.

(Diagram Present in Notes)

- A visual representation of **PaaS** shows how users interact with applications while the cloud provider manages the platform layer.

3. Software as a Service (SaaS)

- Users access fully managed software applications.
- All layers (application, platform, and infrastructure) are **handled by the operator**.
- Users simply **consume the service** without managing any underlying components.

(Diagram Present in Notes)

- A **SaaS diagram** shows how all layers (app, platform, and infrastructure) are managed by the **CSP**.

4. Anything as a Service (XaaS)

- An **extended model** where **any IT service** is delivered as a cloud service.
- Underlying infrastructure **can be managed by either CSPs or third parties**.

Cloud Deployment Models

Cloud deployment models define how cloud infrastructure is set up, who can access it, and its intended use. The key deployment models are:

| Model | Description |
|-----------------|--|
| Public Cloud | Large-scale cloud infrastructure available to the general public. |
| Private Cloud | Cloud infrastructure dedicated to a single organization, offering better control and security. |
| Community Cloud | Shared cloud infrastructure among a specific group of organizations with common interests or policies. |
| Hybrid Cloud | A combination of private and public clouds for enhanced flexibility. |
| Multicloud | Utilization of multiple cloud services from different providers. |
| Federated Cloud | A collaboration of multiple cloud providers, maintaining interoperability. |

Key Comparisons

- **Public Cloud:** Open to everyone, but security can be a concern.
- **Private Cloud:** More control and security but requires investment.
- **Community Cloud:** Restricted to a specific group (e.g., universities, research institutes).
- **Hybrid Cloud:** Best of both public and private clouds.
- **Multicloud:** Helps avoid vendor lock-in by using services from multiple providers.
- **Federated Cloud:** Enables seamless migration across different cloud providers.

Example Implementations:

- **Public Cloud Providers:** AWS, Google Cloud, Microsoft Azure.
- **Private Cloud Examples:** NED.

- **Federated Cloud Use Case:** Used when organizations need cross-cloud credentials and registrations.

NIST Cloud Reference Architecture

Key Components:

1. Stakeholders in the Cloud Ecosystem:

- A. **Users**
- B. **Cloud Service Providers (CSPs)**
- C. **Internet Service Providers (ISPs)**

2. Cloud Broker

- A. Acts as an **interface** offering cloud services on behalf of a cloud operator.
- B. In the past, **cloud brokers were crucial**, but now they are **less common** due to improved **self-service capabilities**.

3. Cloud Broker Roles:

- A. Aggregates services from multiple clouds.
- B. Less relevant today since cloud providers now offer direct self-service.
- C. **Example:** AWS provides its own cloud services without external brokers.

4. Cloud Balancer

- A. Functions as a **service arbitrator**.
- B. Manages communication and resource distribution between cloud providers and users.

5. Cloud Auditor

- A. Conducts audits for **quality assurance, security compliance, and performance benchmarks**.
- B. Auditors must be **independent** and not **internally controlled by the cloud provider** to ensure security and reliability.

6. Cloud Carrier

- A. Acts as an **ISP**, delivering cloud services to end users.

7. Cloud Provider

- A. Responsible for managing **hardware resources and service orchestration**.

Cloud Orchestration & Management

Cloud orchestration involves **automating the deployment, coordination, and management** of cloud resources.

Orchestration Levels:

1. **Cloud Service Provider (CSP) Level**
 - A. Manages large-scale cloud infrastructure.
 - B. Handles automation of cloud resources.
2. **User Level**
 - A. Automates individual functions within the cloud system.

Security & Privacy

- Implemented at **all layers**:
 - **Hardware**
 - **Network**
 - **Software**
 - **User Interfaces**
- Ensures **secure data transmission and access control**.

Cloud Service Management Layers

1. **Business Support**: Provides services and **billing automation**.
2. **Provisioning & Configuration**: Helps in **cloud deployment and workload management**.
3. **Portability**: Enables **migration between cloud providers**.
4. **Interoperability**: Ensures compatibility between different **cloud platforms**.

Advantages & Disadvantages of Cloud Computing

Advantages

1. **Lower Infrastructure & Hardware Costs**
 - A. Reduces the need for upfront capital investment.
 - B. Organizations save money by using **pay-as-you-go** pricing models.

2. Scalability

- A. Cloud resources can expand or contract **on-demand**.

3. Security & Compliance

- A. **Sovereign clouds** offer high security and privacy compliance.
- B. Large-scale cloud operators invest in **robust security measures**.

4. Efficient Resource Utilization

- A. Large-scale cloud providers optimize **performance and storage**.

Disadvantages

1. Handling Scale

- a. Cloud adaptation is growing **exponentially**, leading to **scalability challenges**.

2. Cost Management

- a. Operational costs can **increase over time** if not managed properly.

3. Privacy & Security Concerns

- a. Despite major investments, cloud environments still face **security threats**.

4. Distributed System Challenges

- a. Network latency, system failures, and **data consistency issues** are common in distributed cloud architectures.

Virtualization

1. Virtualization Overview

- Virtualization can be implemented **natively** or **without** cloud computing.
- Most **cloud computing infrastructures rely on virtualization** since it is a **common characteristic** that helps achieve essential cloud features.
- **Purpose of Virtualization:**
 - Efficient utilization of key resources.
 - Provides **abstraction** by modeling hardware or software environments.
 - Enables the creation of **interfaces** that simulate physical components.
 - Utilized by **hypervisors** to create and manage virtual machines.

2. Benefits of Virtualization

- **Resource Management:** Optimizes computing power, storage, and network resources.
- **Ease of Migration:** Virtual machines (VMs) can be migrated easily compared to physical machines.
- **Flexibility:** Virtualization allows multiple applications to run on a single physical resource.
- **Cost Savings:** Reduces the need for physical hardware.

3. Levels of Virtualization

Virtualization can be implemented at different levels in cloud computing:

1. **Server/Machine Virtualization**
 - A. Enables multiple VMs to run on a single **physical machine**.
 - B. Abstracts hardware resources.
2. **Network Virtualization**
 - A. Allows multiple **virtual networks** to operate on shared **physical infrastructure**.
 - B. Example: **Software-defined networking (SDN)**.
3. **I/O Virtualization**
 - A. Virtualizes input/output devices like **storage, network interfaces, and disk controllers**.

4. Virtualization Approaches

Virtualization involves creating **virtual objects** in different ways:

A. Multiplexing

- **Definition:** Creates multiple **virtual objects** from a single **physical resource**.
- **Example: Processor Virtualization** (One CPU appears as multiple virtual processors).

- **OS-Level Multiplexing:** The OS distributes processor time among multiple tasks, giving an illusion of parallel execution.

B. Aggregation

- **Definition:** Combines multiple **physical objects** to create a **single virtual object**.
- **Example: RAID (Redundant Array of Independent Disks)**
 - Combines multiple storage disks into one **logical storage unit**.
 - Provides redundancy for fault tolerance.
- **Diagram Present in Notes:**
 - A simple RAID structure showing multiple disks aggregated into a single logical storage entity.

C. Emulation

- **Definition:** Creates a **virtual object** from a **different type of physical object**.
- **Example: Virtual Memory**
 - The OS **uses disk storage as an extension of RAM**.
 - Allows systems to run applications that require more memory than physically available.

5. Types of Virtualization

A. Hardware-Level Virtualization

- **Large-scale virtualization** used in cloud environments.
- **Type 1 Hypervisor (Bare Metal)**
 - Runs **directly on hardware**.
 - Highly **scalable and efficient**.
 - Examples: **VMware ESXi, KVM, Xen**.
- **Type 2 Hypervisor (Hosted)**
 - Runs **on top of an existing OS**.
 - Less scalable but easier to deploy.
 - Examples: **VirtualBox, VMware Workstation**.
- **Diagram Present in Notes:**
 - Shows a **comparison** between Type 1 and Type 2 hypervisors.

B. System-Level Virtualization

- Allows multiple **operating systems** to run on a **single hardware system**.
- Example: **VirtualBox, VMware Workstation**.
- **Diagram Present in Notes:**
 - Illustrates the relationship between **Guest OS, Host OS, and Hardware** in system-level virtualization.

C. Application-Level Virtualization

- Virtualizes **software applications**, allowing them to run **independently** of the underlying OS.
- Example: **JVM (Java Virtual Machine)** runs Java applications across different platforms.

6. Nested & Memory Virtualization

- **Memory Virtualization** allows the OS to use part of a **physical disk** as virtual memory.
- **Nested Virtualization:**
 - Virtualization **within** virtualization (e.g., running a **Type 2 hypervisor inside a Type 1 hypervisor**).
 - Can degrade performance but is useful for **testing environments**.

7. Network Virtualization

- Inspired by traditional **hardware network functions**.
- Allows **networking components (switches, routers, firewalls, etc.)** to run as **virtual instances**.
- **Infrastructure as a Service (IaaS)** enables users to create virtual networks.
- **Examples:**
 - **VNF (Virtual Network Function)**
 - **NFV (Network Functions Virtualization)**
- **Diagram Present in Notes:**

- Shows multiple **logical networks** running on shared **physical infrastructure**.

Hypervisors

1. Hypervisor Overview

A **hypervisor** is a software layer that allows multiple virtual machines (VMs) to run on a single physical machine by **abstracting hardware resources**.

Responsibilities of a Hypervisor

- **Process Isolation:** Ensures that VM processes do not interfere with each other's memory space (prevents buffer overflow).
- **Resource Management:** Allocates CPU, memory, and storage to VMs efficiently.
- **Virtual Machine Execution:** All VMs are created and managed on top of a hypervisor.
- **Hardware Access Control:** Ensures that instructions requiring hardware access execute under the hypervisor's control.
- **Operating System Modes:**
 - **Privileged Mode:** System-level access.
 - **Non-Privileged Mode:** Application-level access.

2. Key Features of Hypervisors

A. VM Process Isolation (Kernel-Level)

- Prevents VMs from **overlapping or interfering** with each other's memory or processing space.

B. Device Mediation & Access Control

- **Hypervisors emulate hardware devices**, allowing VMs to interact with peripherals.
- Example: When using **VirtualBox** or **VMware Workstation**, pressing a key inside the VM is handled by the hypervisor.

C. Direct Execution of Commands from VMs

- **Privileged Instructions** from guest VMs (e.g., hardware access) must go through the hypervisor.
- **Type 1 Hypervisors** execute commands directly.
- **Type 2 Hypervisors** depend on the base OS for execution.

D. VM Lifecycle Management

- Manages **creation, backup, migration, and deletion** of virtual machines.
- **Live Migration**: Moving a running VM from one server to another without downtime.
- **Snapshot & Backup**: VMs can take periodic snapshots to restore state in case of failures.
- **Data Retention Policies**: Every cloud provider defines policies for VM data storage and backup.

E. Hypervisor Platform Management

- **Updates & Version Control**: Hypervisors receive periodic updates for security and performance improvements.
- **Management Interfaces**:
 - **CLI (Command Line Interface)**: Preferred for large-scale deployments.
 - **Web Interfaces**: Available for ease of use.
- **Security Considerations**:
 - **Type 1 Hypervisors** are more secure since they run directly on hardware.
 - **Type 2 Hypervisors** introduce additional risk due to dependency on the host OS.

3. Types of Hypervisors

Hypervisors are classified based on how they access hardware:

A. Type 1 Hypervisor (Bare Metal)

- Runs **directly on the physical hardware** without an underlying OS.
- **More efficient, secure, and scalable.**
- Examples:
 - **VMware ESXi**
 - **KVM (Kernel-based Virtual Machine)**
 - **Xen**
- **Used in:** Large-scale cloud environments.

B. Type 2 Hypervisor (Hosted)

- Runs **on top of an existing OS**, which manages hardware access.
- **Less efficient and secure** due to OS dependency.
- Examples:
 - **VMware Workstation**
 - **VirtualBox**
- **Used in:** Testing and development environments.

Comparison Diagram in Notes:

- Illustrates **Type 1 Hypervisor running directly on hardware** vs. **Type 2 Hypervisor relying on an OS.**

4. Virtualization Models

There are two primary models for virtualization:

A. Full Virtualization

- **Hypervisor exposes the same hardware interface** to the virtual machine as it does in physical hardware.

- **OS inside VM does not need modification.**
- **Higher overhead** due to full emulation.

B. Para-Virtualization

- **Hypervisor provides a specialized interface** that does not exist in physical hardware.
- **Guest OS must be modified** to communicate with the hypervisor.
- **Performance optimization** due to lower overhead.

Comparison Table in Notes:

| Feature | Full Virtualization | Para-Virtualization |
|--------------------|------------------------|---|
| Hardware Interface | Same as physical | Optimized for VM |
| OS Modification | Not required | Required |
| Performance | Lower due to emulation | Higher due to direct hypervisor interaction |

5. Hypervisors Used in Cloud Computing

- **Xen Hypervisor:**
 - Used by **AWS** (with modifications).
 - **AWS Nitro Hypervisor** is based on Xen.
 - Supports both **para-virtualization and full virtualization**.
 - Used in **ARM processors** for mobile and embedded systems.
- **KVM (Kernel-based Virtual Machine)**
 - Open-source Linux hypervisor.
 - **Technically a Type 1 hypervisor** but runs alongside the Linux OS.
 - Allows direct access to hardware.
 - **Default hypervisor in Linux-based cloud environments.**
- **VMware ESXi**
 - A **commercial Type 1 hypervisor** developed by VMware.
 - **Bare-metal installation** for cloud computing.
 - Expensive licensing but widely used in enterprise environments.

Diagram Present in Notes:

- Shows **KVM architecture**, depicting **VMs running on a Linux kernel hypervisor**.

6. SWOT Analysis of Hypervisors

A SWOT (Strengths, Weaknesses, Opportunities, Threats) analysis was conducted to assess hypervisors' role in cloud computing.

Strengths

- **High scalability** with Type 1 hypervisors.
- **Efficient resource allocation** and isolation.
- **Support for live migration and snapshots**.

Weaknesses

- **Type 2 hypervisors depend on the host OS**, making them less efficient.
- **Security vulnerabilities in Type 2 hypervisors** due to OS dependency.

Opportunities

- **Increased adoption of open-source hypervisors** (e.g., KVM) in cloud computing.
- **Integration with AI and automation** for smarter VM management.

Threats

- **Cybersecurity risks** in cloud environments.
- **Vendor lock-in** when using proprietary hypervisors like VMware ESXi.

Containers & Kubernetes

1. Containers

Overview

- **Inspired by application-level virtualization.**
- **Difference from Virtual Machines (VMs):**
 - **VMs** contain a full-fledged OS and take up several GBs.
 - **Containers** are lightweight, providing an **isolated environment** to run specific applications.
- **Diagram in Notes:**
 - Shows a **comparison of containers and VMs**, with VMs running on a hypervisor and containers running on a container engine.

Advantages of Containers

- **Smaller Footprint:** Consumes fewer resources than traditional VMs.
- **Faster Execution:** Containers spawn quickly compared to VMs.
- **Efficiency:** Instead of running an entire OS, containers **share the host OS kernel**.

Containerization Evolution

- **Previously:** Workloads were executed using **Virtual Machines**.
- **Now:** Workloads are executed in **Containers**.
- **Popular Container Platforms:**
 - **Docker:** Leading containerization technology.
 - **Open Container Initiative (OCI):** Standardization of container runtimes.

Container Optimization

- Optimization depends on the **container engine**, which impacts:
 - **Performance**
 - **Security**
 - **Scalability**
- **Orchestration Tools:**
 - **Kubernetes**
 - **Docker Swarm**
- **Elasticity:**

- Containers are **more elastic** than virtual machines.
- They can be created and destroyed rapidly.

Challenges with Containers

- **Security Issues:**
 - Containers have a **short lifespan**, making security auditing and investigations difficult.
- **Tracking Issues:**
 - Short-lived containers make **monitoring and logging** challenging.

2. Cloud-Native Applications (CNA)

Definition

- Applications designed **specifically for cloud environments**.
- Built using **containers and microservices**.

Cloud-Native Characteristics

- **Highly Scalable**
- **Flexible**
- **Fault-Tolerant**
- **Lightweight & Efficient**
- **Supports Continuous Deployment**

Microservices Architecture

- **Modular Approach:** Applications are split into **small, independent services**.
- **Cloud-Optimized:** Unlike monolithic applications, microservices leverage cloud elasticity.
- **Examples:**
 - **YouTube**
 - **Netflix**
 - **AWS Cloud Services**

Diagram in Notes:

- Illustrates **Cloud-Native Architecture**, showing the breakdown into **microservices, containers, and Kubernetes**.

3. Kubernetes

Overview

- Kubernetes is a **container orchestration platform** that **manages and scales containers** across multiple machines.
- **Inspired by Google's internal Borg system.**

Kubernetes Architecture

- **Control Plane:**
 - Manages and schedules workloads across worker nodes.
- **Worker Nodes (Computing Plane):**
 - Run the actual containerized workloads.

Diagram in Notes:

- Shows a **Kubernetes cluster** with:
 - **Control Plane**
 - **Worker Nodes**
 - **Pods running containers**

Key Kubernetes Components

1. **Node:**
 - A. A machine (physical or virtual) where workloads run.
2. **Pod:**
 - A. Smallest deployable unit in Kubernetes.
 - B. Can contain **one or more containers**.
3. **Scheduler:**
 - A. Distributes workloads across nodes.
4. **etcd:**
 - A. Persistent storage (key-value store for configuration data).
5. **API Server:**

- A. Manages API requests.
- 6. **Controller Manager:**
 - A. Maintains the state of Kubernetes objects.
- 7. **Kube-proxy:**
 - A. Manages network communication between pods.

Diagram in Notes:

- Depicts the **Kubernetes Control Plane and Components**.

4. Kubernetes & Docker

Docker vs. Kubernetes

- **Docker:**
 - Creates and runs containers.
- **Kubernetes:**
 - Manages, orchestrates, and scales containers across clusters.

Kubernetes Managed Services

- **Amazon Elastic Kubernetes Service (EKS)**
- **Azure Kubernetes Service (AKS)**
- **Google Kubernetes Engine (GKE)**

Kubernetes Deployment on Cloud

Two Approaches

1. **System Build (Manual Deployment)**
 - A. Build nodes (master, slave/worker).
 - B. Not auto-scalable.
2. **Kubernetes as a Service (Managed)**
 - A. Fully managed service that provides:
 - i. **Cluster provisioning**
 - ii. **Elastic scaling**
 - iii. **Auto-management**

Kubernetes Scalability

- Kubernetes **adds worker nodes dynamically** as needed.
- In research-based tasks, **Kubernetes as a Service may not be suitable** because it bypasses certain configurations.

Security & Logging Considerations

- Since containers **spawn and terminate quickly**, **logging is crucial** for tracking container events.
- **Red Hat Kubernetes** provides **enterprise-level** Kubernetes solutions.

AWS & Virtual Private Cloud (VPC)

1. Amazon Web Services (AWS)

Overview

- **AWS** is the **largest public cloud service provider**.
- Provides **Infrastructure as a Service (IaaS)**, allowing users to build **entire cloud-based infrastructure**.

Key Benefits of AWS

1. **Security** – Ensures data protection and access control.
2. **Availability** – Highly redundant infrastructure with global coverage.
3. **Performance** – Optimized cloud performance with scalability.
4. **Scalability** – Allows seamless growth and resource provisioning.
5. **Flexibility** – Multiple configurations for different workloads.
6. **Global Footprint** – Data centers distributed worldwide.

Security & Availability Considerations

- Security is the **top priority** since AWS is a **public cloud**.
- **Availability** is an attribute of security but is **listed separately** because:

- The region where a service is available impacts **privacy and performance**.
- **Latency is a key benchmark** for IoT-based applications on AWS.

Edge & Fog Computing in AWS

- **Edge Computing** reduces latency by processing data closer to users.
- **Fog Computing** acts as a middle layer between cloud servers and edge devices, improving real-time processing.
- **Used for IoT applications** requiring real-time processing.

Scalability & Resource Management

- **Over-provisioning** ensures resources are always available when needed.
- Scalability **has limits**, making it a **classic challenge in cloud computing**.
- **Admission Control** enforces **resource limits per user**.
- **Flexibility** in AWS includes choosing **data center locations** and **storage options**.

Global Infrastructure of AWS

- AWS **data centers** are typically located near **sea ports** for high-speed **fiber-optic connectivity**.
- **Networking backbone** consists of **Wide Area Network (WAN)** and **Local Area Network (LAN)** for **efficient data transfer**.

2. Virtual Private Cloud (VPC)

What is a VPC?

- **VPC (Virtual Private Cloud)** is a **logically isolated section** within AWS, providing **dedicated cloud space for clients**.
- **Multi-Tenant Architecture**:
 - Multiple users can use AWS, each assigned **their own VPC**.
 - Users **can have multiple VPCs** based on **billing and workload requirements**.

Diagram Present in Notes:

- Shows a **VPC with a VM** inside it, demonstrating logical isolation.

VPC Hierarchy

1. **VPC → Virtual Machines (VMs) → Operating System (OS) → Platform → Applications**
2. **For IaaS:** The cloud provider **allocates a VPC** and users build their infrastructure.
3. **For PaaS/SaaS:** The stack is managed by the cloud provider.

Logical Isolation in AWS

- **Hypervisors manage VM isolation.**
- **Network Policies & Access Controls** ensure isolation within a VPC.
- **Logical isolation** is enforced, but **physical isolation is difficult.**
- **VPCs from different clients must be isolated.**
- **Same IP addresses can be used by different clients** since they are only valid within their **VPC scope.**

3. AWS Compute Services: Elastic Compute Cloud (EC2)

Overview

- **AWS EC2** provides virtual machines (VMs) in the cloud with **on-demand resource allocation.**
- **Features:**
 - **VM Creation & Management** – Automated provisioning.
 - **Compute Power (CPU, RAM, Storage)** – Configurable based on workload.
 - **Pre-Built OS Images** – OS is already installed and can be attached to VMs.
 - **Network-Based Storage** – Uses AWS **Storage Servers** for backups.

Diagram Present in Notes:

- Shows **two VMs with OS and storage connected to a network-based storage server.**

EC2 Variants

1. **Standard EC2** – General-purpose virtual machines.
2. **GPU-Enabled EC2** – Specialized VMs with **dedicated GPU resources**.
3. **Dedicated EC2** – Provides **physical isolation**, bypassing the hypervisor.
4. **HPC EC2** – High-performance computing for **massive parallel workloads**.

AWS EC2 Pricing Models

- **Auctioned EC2 Instances:**
 - Allows **bidding on spare EC2 capacity**.
 - **Cost-effective** for non-critical workloads.
 - **Highest bid wins** the instance.
- **Billing Variations:**
 - **Pricing varies by region** due to energy and infrastructure costs.
 - **Different services have different pricing** based on the region.

EC2 Deployment Considerations

- **Users specify where their EC2 machine is created** (region selection).
- **Operators decide which physical server** to assign within that region.
- **Schedulers manage** the resource allocation from a pool of servers.

AWS Infrastructure Services

When a machine is created, AWS handles:

1. **Storage Data Placement** – Decides the exact physical server for data.
2. **Compute Services (EC2)**
3. **Network Services** – Assigns **IP addresses**.
4. **Identity & Access Management (IAM)** – Manages user roles and authentication.
5. **Storage Services** – Provides OS images and backup options.

Diagram Present in Notes:

- Shows **AWS infrastructure flow**, linking **IAM, EC2, Network, and Storage Services**.

AWS Service Dependencies

- Some services are **critical** for machine deployment.
- **A machine cannot be created without:**
 - **OS (Storage Service)**
 - **Network Services (IP Address Assignment)**

1. Elastic IP Address

Overview

- **Elastic IP (EIP)** is a **static IP address** allocated from **AWS's pool of public IPs**.
- It is **not tied to a single EC2 instance** but can be reassigned.
- When a machine is created, it receives an **IP assigned via DHCP**, which **changes over time** unless an **Elastic IP** is assigned.

Elastic IP Functionality

- **Used when a fixed IP address is required**, such as for **web hosting or load balancing**.
- **Diagram Present in Notes:**
 - Shows **EC2 instances with dynamic IPs** connected to a **load balancer** with an **Elastic IP assigned to the account**.

Key Points

- Web servers (EC2) have **dynamically assigned IPs**.
- Load balancers distribute traffic and manage the flow between **users and EC2 instances**.
- Elastic IP is **static** and remains fixed even if an instance is restarted.
- **Auto-scaling ensures availability**, creating new instances as needed.

2. Database Services

Overview

- AWS provides **multiple database solutions**, including **managed and self-hosted options**.
- **Key Database Types:**
 - **Relational Databases (RDS)** – Structured storage with SQL support.
 - **Non-Relational Databases (NoSQL)** – Schema-less, scalable databases.
 - **Database Migration Services** – Used for migrating data between databases.

3. Storage Services

Types of File Storage in AWS

1. **Network-Based Storage:**
 - A. Distributed File System (DFS) ensures scalability and redundancy.
 - B. Multiple storage options exist in **AWS VPC**.
2. **Object Storage (S3):**
 - A. Stores data as objects rather than files.
 - B. **OS cannot be installed on S3**, as it is used for **storing backups, media, and application data**.
 - C. **Used as Platform as a Service (PaaS)**.
 - D. Capacity, access speed, and **pricing models vary**.
3. **Block Storage (EBS):**
 - A. Used for persistent storage on **EC2 instances**.
 - B. Supports **OS installation**.
 - C. Processing to retrieve blocks can be **slower compared to direct disk access**.
4. **Archival Storage (Glacier):**
 - A. Designed for **long-term data retention**.
 - B. **Very low cost**, but **high latency** for retrieval.

Diagram Present in Notes:

- Shows **different AWS storage types, including Object, Block, and Archival storage**.

4. AWS Networking

Overview

- AWS provides a **robust networking infrastructure** to connect and manage cloud resources.
- Networking in AWS consists of:
 - **Backbone Networks** – Data center interconnectivity using **fiber-optic infrastructure**.
 - **Virtual Networks** – User-defined networks within AWS.

Key Components

- **VPN (Virtual Private Network):**
 - Extends **on-premises networks** to AWS securely.
 - Makes users part of a **private cloud** setup.
- **VPC (Virtual Private Cloud):**
 - Provides **logical network isolation** for AWS users.
 - **Allows multiple VMs to run within a secured environment.**
- **Virtual Routers & Subnets:**
 - AWS **Virtual Router (V. Router)** handles internal network traffic.
 - **Subnetting** allows segmentation within a VPC.

Diagram Present in Notes:

- Shows **VPC with connected VMs, Virtual Routers, and external networking infrastructure.**

5. AWS Deployment & Administration

Automation in AWS

- Infrastructure **can be automated** for deployment.
- **Key AWS Automation Tools:**
 - **AWS CloudFormation** – Infrastructure as Code (IaC).

- **AWS OpsWorks (Chef-based automation)** – Configuration management.

Monitoring Services

- **AWS CloudTrail** – Tracks **API requests** and logs actions.
- **AWS CloudWatch** – **Real-time monitoring** of AWS resources.

Service Level Agreements (SLAs)

- Defines the **level of service** between **AWS and the user**.
- Responsibilities differ **based on service type (IaaS, PaaS, SaaS)**.

6. AWS Networking Firewalls

Overview

- AWS **blocks all external access** by default.
- **Firewalls are necessary** for EC2 instances and other AWS resources.

Key Components

- **Firewall Policies:**
 - **Users control firewall rules** via **AWS Security Groups** and **Network Access Control Lists (NACLs)**.
- **SSH Security:**
 - Transport Layer Security (TLS) ensures **secure access to AWS resources**.
 - **Users are responsible** for setting up firewall rules.

AWS Content Delivery Networks (CDN)

- AWS **CloudFront** speeds up **content delivery globally**.
- **CDN provides region-based content delivery** for performance and compliance.

7. Deployment Models in AWS

Key Points

- Deployment can be **automated or manual** (via scripts).
- Multiple service configurations exist:
 - **2-Tier & 3-Tier Architectures.**
 - **Customized service models** based on use case.

Performance & Reliability Considerations

- **System performance and reliability** are major concerns for cloud deployment.
- **High reliability can only be provided by large cloud service providers (CSPs).**

OpenStack

OpenStack is an **ecosystem** that allows for the construction of cloud environments using open-source software.

Cloud Construction Types

Using OpenStack, multiple types of cloud environments can be created:

- **Public Cloud**
- **On-premises Cloud (Private Cloud)**
- **Edge Cloud** (Created at the edge for IoT, Telecom systems)

Key Features

- **Open-source and free**
- **Commercialized versions** also exist
- **Apache Cloud Stack** serves as an alternative to OpenStack, along with **Open Nebula**.
- **Base metal, virtual machines, and containers** can be handled via OpenStack.

OpenStack and AWS

- AWS provides **bare metal as EC2 instances** along with other services similar to OpenStack.
- **The world runs on OpenStack:**
 - **Case Studies:** Exploring organizations using OpenStack.
 - **OpenInfra:** A project of OpenStack.
 - **Contributors:** NASA, Rackspace Cloud.
 - **Use cases** of OpenStack across industries.

OpenStack Deployment

- A cloud can be built on OpenStack, and additional services and features can be integrated.
- Deployment can be **automated** as discussed in prior sections.
- **OpenStack is primarily an IaaS provider.**
- OpenStack can be:
 - Downloaded as **open-source software**.
 - Acquired as **vendor-specific OpenStack solutions** (e.g., **RedHat OpenStack, Ubuntu OpenStack**).
- **Installation Requirements:**
 - **Multiple Machines** or
 - **DevStack (Single Machine Setup):** A script used to construct a cloud on a **laptop** (usually not feasible for production but useful for **testing purposes**).

OpenStack Architecture

1. Dashboard Service

- Provides a **web interface** for managing OpenStack.
- Apache-based, designed for **end-users and administrators** to access control.

2. Compute Service

- Manages **VM lifecycle**.
- **Hypervisor:** **KVM** (default, but AWS uses different hypervisors).

- **Responsibilities:**
 - Distributes **resources across machines**.
 - Manages **VM scheduling decisions**.
- OpenStack's **compute service can be modified**, but AWS compute services cannot.

3. Networking Service

- Manages networking functionalities:
 - Firewall rules
 - IP assignment
 - VPN routing
- Uses **Open vSwitch & SDN controllers** for **DHCP, routing, and network firewall management**.

4. Storage Service

- Manages **distributed file systems**.
- Can be used as:
 - **A network service**
 - **An object service**
- Used to **build storage-as-a-service** in the cloud.
- **Supports multiple types of storage:**
 - **Block storage**
 - **Object storage** (default offerings).

5. Identity Service

- Provides **authentication and authorization** for OpenStack services and users.
- **Ensures isolation** in IaaS cloud:
 - **Access Control**
 - **Compute Service Isolation (Hypervisor-based)**
 - **Network-Level Isolation**
- Identity services can be modified for security enhancements.

6. Image Service

- Handles **VM image creation and management**.

- Similar to **Amazon Machine Images (AMIs)**.

7. Orchestration Service

- Automates the **management and scaling** of multiple machines.

8. Monitoring/Telemetry Service

- Collects metrics and performance data for **monitoring OpenStack services**.

9. Database Services

- Supports:
 - **Relational Databases**
 - **Non-Relational Databases**
 - **Migration between databases**

Virtual Machine Lifecycle in OpenStack

- When creating a **VM**, several services are involved:
 - **Telemetry/Monitoring Dashboard**
 - **Identity Service**
 - **Compute Service**
 - **Image Service**
 - **Networking Service**
 - **Storage Orchestration**
- These services operate on a **distributed system** and must be coordinated properly.
- **Issues that can occur:**
 - **Network service failure** → Machine is created but has **no IP**.
 - **Orchestration & Dashboard are optional**, but identity and image services **are necessary**.
 - **Storage service failure** → Machine is created, but **no storage** is allocated.

Master-Slave Architecture in OpenStack

- **Cloud Controller** → **Master Node** (Component Management)
- **Compute Nodes, Network Nodes, Storage Nodes** function under master-slave architecture.

OpenStack Node Architectures

Three-Tier Architecture of OpenStack

1. **Controller Node**
 - a. Manages overall OpenStack services.
 - b. Handles authentication, APIs, and orchestration.
 - c. Responsible for monitoring and scheduling.
2. **Compute Node**
 - a. Runs virtual machines (VMs) using a hypervisor (e.g., KVM).
 - b. Manages VM execution and resource allocation.
3. **Networking Node**
 - a. Handles network connectivity for OpenStack.
 - b. Manages virtual networking, routing, and IP assignments.

Note: A minimum of three separate machines is required for this architecture.

Two-Tier Architecture of OpenStack

1. **Controller Node**
 - a. Combines API, identity services, orchestration, and monitoring.
2. **Compute + Networking Node**
 - a. A single node manages both virtual machines and networking.

Example: If **3,000 machines** are used to construct a cloud, **90% of them** are utilized for compute nodes.

- **Storage as a Service Cloud:**
 - Most machines are allocated for storage.
- **DevStack:**
 - Consolidates all nodes into a **single physical machine**.
 - Used **only for testing** due to its **limited scalability**.

Conceptual vs. Physical Architecture

- **Conceptual/Logical Architecture**
 - Defines how a VM is launched.
 - Identifies services involved.
 - Focuses on distributed services.
- **Physical Architecture**
 - **Single Node:** Minimum of **1 Controller**.
 - **Two Nodes:** Minimum of **1 Controller + Compute/Networking**.
 - Defines roles of machines and how resources are allocated.

Components of OpenStack

1. **SQL Database Service**
 - a. Stores metadata of the cloud (logs, services).
 - b. **Not** a Database-as-a-Service (DBaaS).
2. **Message Queue Service**
 - a. Ensures services **communicate via messaging**.
3. **Network Time Service**
 - a. Synchronizes machine time.
 - b. Uses **GMU** (Global Master Unit) to ensure all nodes share the same clock.
 - c. The **Controller Node** acts as the **time source**.
4. **Identity Service**
 - a. Provides **authentication** and **authorization**.
5. **Networking ML2 Plugin**
 - a. Manages **Layer 2 networking** services.

Compute Nodes

- **KVM Hypervisor**
 - Default hypervisor used in OpenStack.
- **Open vSwitch**
 - Allows **unlimited VMs** to connect.
 - Software-based **network switch** used for networking and compute.

Network Nodes

- **L3 Agent**
 - Connects cloud machines to the external world.
 - Provides **routing**.
- **DHCP Agent**
 - Assigns **IP addresses** dynamically.
- **Open vSwitch**
 - Enables software-defined networking (SDN) in OpenStack.

Note: Many networking and compute components are merged in **small-scale architectures**.

Storage Service

- **Mandatory for OpenStack** (not optional).
- **Telemetry Agent**
 - Manages **billing, resource monitoring, and prediction**.
 - **Collects stats** from OpenStack nodes.

Privacy & Security in OpenStack

- **Red Interface (Network Management)**
 - Manages internal communication.
- **Yellow Interface (Compute Node Communication)**

- Used for **fast data transfers** when a new machine is created.
- **Separate Storage Network**
 - Enables **high-speed data transfer** between compute nodes and storage servers.

External Networking in OpenStack

- **External Network**
 - Provides routing for **connections to the external world**.
 - Worker and compute nodes are **not exposed** by default for security reasons.
- **Logical vs. Physical Network**
 - **Logical Network:** Manages all resources (SDN-based).
 - **Physical Network:** Hardware-level networking.

Diagrams Present in Notes:

- **Storage Network Diagram**
 - Shows interaction between **VMs, network storage, and compute nodes**.
- **External Network Diagram**
 - Explains how OpenStack nodes connect to **external services** securely.